

```
In [56]: import pandas as pd
data= pd.read_csv('https://raw.githubusercontent.com/SivadineshPonrajan/Machine-Learning-Notebooks/master/Dataset/Churn_Modelling.csv')
data.head()
```

Out[56]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1

```
In [57]: process=data.iloc[ :, [3,4,6,7,8,9,10,11,12,13] ]
process.head()
```

Out[57]:

	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	France	42	2	0.00		1	1	101348.88	1
1	608	Spain	41	1	83807.86		1	0	112542.58	0
2	502	France	42	8	159660.80		3	1	113931.57	1
3	699	France	39	1	0.00		2	0	93826.63	0
4	850	Spain	43	2	125510.82		1	1	79084.10	0

```
In [58]: import numpy as np
for i in range(len(process.iloc[:,1])):
    if(process.iloc[:,1][i]=="France"):
        process.iloc[:,1][i]= np.int64(1)
    elif(process.iloc[:,1][i]=="Spain"):
        process.iloc[:,1][i]=np.int64(2)
    else:
        process.iloc[:,1][i]=np.int64(3)
process.head()
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:4: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
after removing the cwd from sys.path.

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:6: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:8: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[58]:

	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	619	1	42	2	0.00		1	1	101348.88	1
1	608	2	41	1	83807.86		1	0	112542.58	0
2	502	1	42	8	159660.80		3	1	113931.57	1
3	699	1	39	1	0.00		2	0	93826.63	0
4	850	2	43	2	125510.82		1	1	79084.10	0

```
In [ ]: x,y=process.iloc[:, :-1],process.iloc[:, -1]
```

```
In [ ]: from sklearn import preprocessing
x = preprocessing.StandardScaler().fit(x).transform(x.astype(float))
x_train,x_test,y_train,y_test=x[:8000:],x[8000:],y[:8000:],y[8000:]
```

```
In [61]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(8000, 9)
(2000, 9)
(8000,)
(2000,)

```
In [62]: from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier(random_state=50,average=True)
sgd_clf.fit(x_train,y_train)
```

Out[62]: SGDClassifier(alpha=0.0001, average=True, class_weight=None, early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True, l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2', power_t=0.5, random_state=50, shuffle=True, tol=0.001, validation_fraction=0.1, verbose=0, warm_start=False)

```
In [69]: print(y_test[30:40:])
sgd_clf.predict(x_test[30:40:])
```

8030 1
8031 0
8032 0
8033 0
8034 0
8035 0
8036 0
8037 0
8038 0
8039 0
Name: Exited, dtype: int64

Out[69]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```
In [64]: from sklearn import metrics
print("Train set acc: " ,(metrics.accuracy_score(y_train, sgd_clf.predict(x_train))))
print("Test set acc: " ,(metrics.accuracy_score(y_test, sgd_clf.predict(x_test))))
```

Train set acc: 0.80375
Test set acc: 0.8105

```
In [73]: training_score = sgd_clf.score(x_train, y_train)
testing_score = sgd_clf.score(x_test, y_test)

print("training_score: "+str(training_score))
y_pred= sgd_clf.predict(x_test)

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix: For Training ")
print(cm)

print("testing_score: "+str(testing_score))
y_pred=sgd_clf.predict(x_train)

cm = confusion_matrix(y_train, y_pred)
print("Confusion Matrix: For Testing ")
print(cm)

training_score: 0.80375
Confusion Matrix: For Training
[[1590 20]
 [ 359 31]]
testing_score: 0.8105
Confusion Matrix: For Testing
[[6280 73]
 [1497 150]]
```

```
In [75]: from sklearn.metrics import classification_report
cr = classification_report(y_train, y_pred)

print("classification Report:")
print(cr)
```

classification Report:
precision recall f1-score support
0 0.81 0.99 0.89 6353
1 0.67 0.09 0.16 1647
accuracy 0.80 8000
macro avg 0.74 0.54 0.52 8000
weighted avg 0.78 0.80 0.74 8000

```
In [ ]:
```